

Introduction to Red Hat Enterprise Linux AI 1.1 on Dell PowerEdge GPU Enabled Servers

Enabled by Dell PowerEdge Servers

September 2024

H20138

White Paper

Abstract

This white paper provides an overview of the features and capabilities of Red Hat Enterprise Linux AI on Dell GPU Enabled Servers.

Dell Technologies AI Solutions

Dell

Reference Design

Copyright

The information in this publication is provided as is. Dell Inc. makes no representations or warranties of any kind with respect to the information in this publication, and specifically disclaims implied warranties of merchantability or fitness for a particular purpose.

Use, copying, and distribution of any software described in this publication requires an applicable software license.

Copyright © 2024 Dell Inc. or its subsidiaries. Published in the USA September 2024 H20138.

Dell Inc. believes the information in this document is accurate as of its publication date. The information is subject to change without notice.

Contents

Executive summary 4

Solution overview 5

Solution design 6

Security considerations..... 9

Conclusion..... 9

References 10

Executive summary

Overview

Red Hat Enterprise Linux AI (RHEL AI) serves as a foundation model platform, seamlessly facilitating the development, testing, training, and inferencing of generative AI (GenAI) models for enterprise applications.

RHEL AI brings together:

- The Granite family of open source Apache 2.0-licensed large language models (LLMs) with complete transparency on training datasets.
- InstructLab model alignment tools, which open the world of community developed LLMs to a wide range of users, that offer fine-tuning of models via the Large-scale Alignment for ChatBots (LAB) technique.
- A bootable image of RHEL including popular AI libraries such as PyTorch and hardware optimized inference for NVIDIA, Intel, and AMD.
- Enterprise-grade technical support and model intellectual property indemnification provided by Red Hat.

RHEL AI allows portability across hybrid cloud environments and makes it possible to scale your AI workflows either on Red Hat OpenShift AI or inference them on RHEL AI.

This white paper presents a reference implementation for RHEL AI on Dell PowerEdge servers. This paper also provides a brief overview of Red Hat Enterprise.

Audience

This document is intended for data scientists, data engineers, AI developers, developers, system administrators, and architects. Experience with RHEL is recommended, but not required.

Revisions

Date	Part number/ revision	Description
September 2024	H20138	Initial release

We value your feedback

Dell Technologies and the authors of this document welcome your feedback on this document. Contact the Dell Technologies team by [email](#).

Contributors: Jacob Kirby (Dell), Deepak Aakula (Dell), Matthew Miller (Red Hat), Jeremy Eder (Red Hat)

Note: For links to other documentation for this topic, see the [Specialty Servers Info Hub](#).

Solution overview

Business challenges

- **Resource Constraints:** Many organizations lack the necessary hardware and user-generated data resources to undertake resource-intensive training or fine-tuning of LLMs. RHEL AI provides a comprehensive suite of fine-tuning utilities (based on LAB) that minimizes dependence on costly human annotations and proprietary models, while allowing simplified adoption and updates using a composed appliance based model that can run on a single server.
- **Freedom and Flexibility:** Allows organizations to quickly get started with generative AI by bringing a solution closer to the data assets, with tooling that is adoptable by both professional data scientists and developers.
- **Compliance and Regulatory Issues:** Many countries have laws that require data to be stored within their borders, thus limiting their capability to use distributed services that are based in other countries. RHEL AI provides a way to fine-tune models within the borders/boundaries of your controlled environment.
- **Security:** Many organizations worry about the security and control of their data. RHEL AI uses Red Hat Enterprise Linux Image Model, which allows users to deploy and manage RHEL AI capabilities with a secure bootc container image. For more information see: <https://www.redhat.com/en/technologies/linux-platforms/enterprise-linux/image-mode>.




Figure 1. RHEL AI Appliance

Business benefits

- **Open Source Innovation:** RHEL AI harnesses open-source models and tools, enabling businesses to capitalize on community-driven innovation and collaboration.

- **Optimized Performance:** The platform incorporates a streamlined system image that benefits from hardware acceleration, ensuring optimal performance across diverse hardware platforms.
- **Enterprise Support:** RHEL AI offers enterprise-grade technical support, ensuring that businesses have the assistance they need to deploy and maintain AI solutions effectively.

Solution benefits




Figure 2. RHEL AI benefits

Solution design

Technology overview

RHEL AI

Includes the following hardware and software components:

- Dell PowerEdge Servers
- Industry-leading NVIDIA GPUs for AI/ML
- Red Hat RHEL AI bootable container image based on RHEL image mode

Dell PowerEdge servers with GPU

RHEL AI has been validated on the following servers:

- PowerEdge XE9680 with 8 x Nvidia H100 SXM GPU
- PowerEdge 760xa with 4 x Nvidia H100 PCIe GPU with NVLink
- PowerEdge 760xa with 4 x Nvidia L40S PCIe GPU with NVLink



Figure 3. PowerEdge XE9680 Server



Figure 4. PowerEdge R760xa Server

NVIDIA GPUs for AI/ML

- Nvidia H100 SXM GPU
- Nvidia H100 PCIe GPU
- Nvidia L40S GPU

Red Hat RHEL AI

RHEL AI is a foundation model platform that combines open-source LLMs, enhanced tooling, and a community-driven approach to generative AI model development.

Deployment models

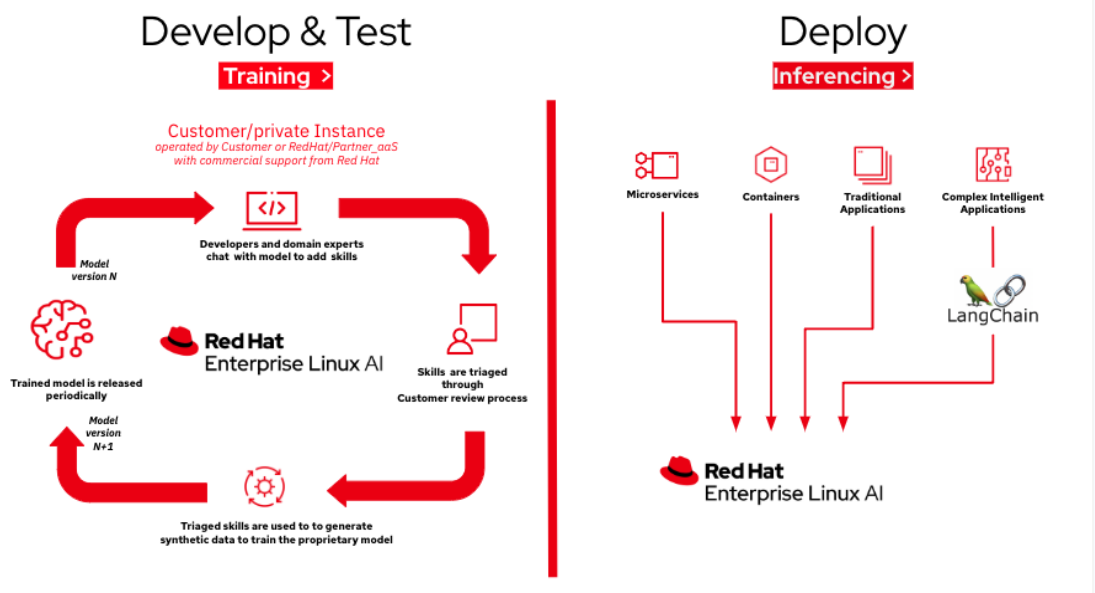


Figure 5. AI Lifecycle using Red Hat Enterprise Linux AI

Hardware design The RHEL AI server is self-contained in a single PowerEdge server including the GPUs(s).

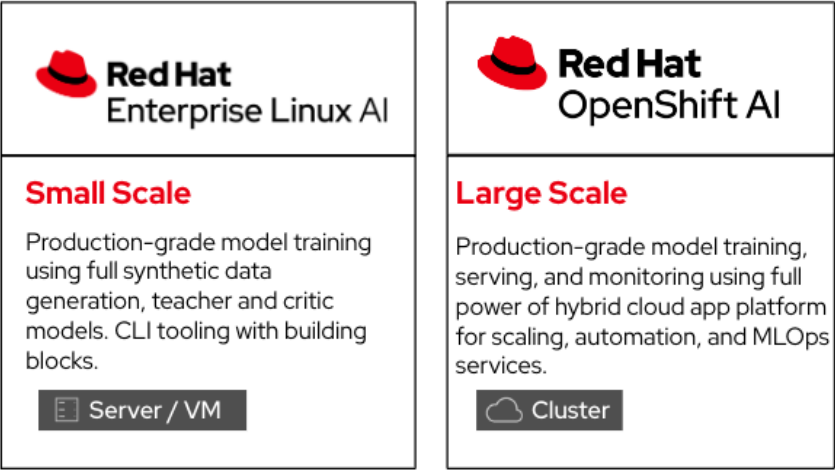


Figure 6. Small and large scale Red Hat Deployment options

Storage RHEL AI leverages the local storage of the PowerEdge server on which it runs. Local disks are recommended to be solid-state storage.

Networking Each RHEL AI server is a standalone server, requiring no specialized networking beyond standard network access for client access, and access to external storage and other AI infrastructure.

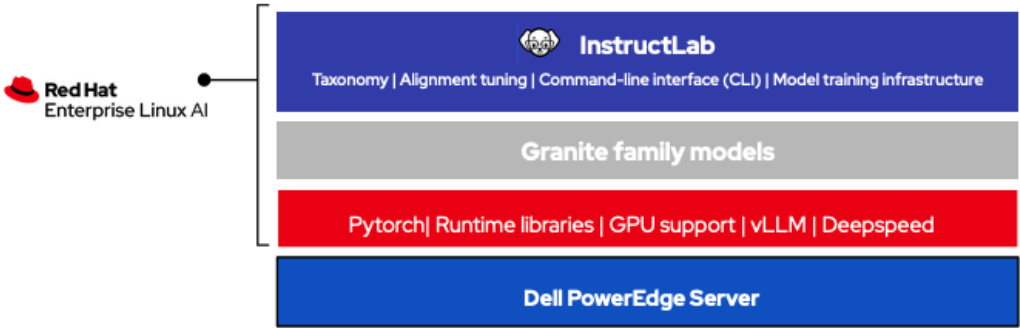


Figure 7. Red Hat Enterprise Linux AI software stack

**Open-source
IBM Granite
models**

The [IBM Granite](#) family of artificial intelligence (AI) models include:

InstructLab

InstructLab, an open-source project created by IBM Research and Red Hat, based on the aforementioned LAB method, aims to enhance LLMs used in generative AI applications by making fine-tuning of foundation models straight-forward and accessible.

It leverages three components:

- **Taxonomy-based data representation:** Humans add and curate knowledge and skill data to the taxonomy for model training.
- **Large-scale synthetic data generation:** A teacher model creates many new examples based on the manually input seed training data. A critic model filters the synthetic data for accuracy and safety.
- **Iterative alignment tuning:** The LLM is retrained using synthetic data.

[InstructLab](#) provides a cost-effective way to enhance pre-trained large language models to fit specific business requirements.

RHEL in image mode

Image mode for RHEL offers a simplified approach to managing the RHEL operating system. It allows administrators to build a bootable container image in the same manner that a container image is built. Container images are also used for image updates. This brings operating system builds, updates, and maintenance into the same CI/CD operational workflows as an application.

Implementation guidance

Refer to Red Hat installation documentation [here](#).

Security considerations

Deploy and use the hardware and software recommended in this solution securely. Follow recommendations from Dell Technologies and the other vendors cited in this technical white paper. For additional information, refer to the Dell Technologies Security and Trust Center at <https://www.dell.com/en-us/dt/about-us/security-and-trust-center/index.htm>.

Conclusion

RHEL is the world's leading enterprise Linux platform, certified on hundreds of clouds and with thousands of hardware and software vendors. With the technological foundation of Linux, containers, and automation, Red Hat's open hybrid cloud strategy gives you the flexibility to run your AI applications anywhere you need them.

RHEL AI, Granite, and InstructLab further deliver on this vision, breaking down the cost and resource barriers to experimenting with and building AI models while providing the tools, data, and concepts needed to fuel the next wave of intelligent workloads.

References

Dell Technologies documentation

The following links provide additional information about the solution design and components used in this design. Other reference links are provided in context in the body of the document.

- [Dell Technologies Info Hub for AI Solutions](#)

Red Hat Enterprise Linux AI documentation

- [Red Hat Enterprise Linux AI documentation](#)
- [Red Hat Enterprise Linux AI | Red Hat Developer](#) - Overview and resources for RHEL AI
- [What is RHEL AI? A guide to the open source way for doing AI](#) - Detailed guide on RHEL AI and its open-source approach
- [Open Source Powered AI/ML for the Hybrid Cloud | Red Hat Developer](#) - Information on AI/ML tools and resources for hybrid cloud environments
- [RHEL AI - GitHub](#) - GitHub repository for RHEL AI projects

Other resources

- [arxiv logo > cs > arXiv:2403.0108: LAB: Large-Scale Alignment for ChatBots](#)
- [Hugging Face > InstructLab](#)
- [InstructLab](#) - A new community-based approach to build truly open-source LLMs